

Payouts API

High-Level Design

Boku — Banking & Settlement

Senior Backend Engineer take-home · 2026

The Task

Design a Payouts API that lets a caller submit a payout to a beneficiary, exposes clear endpoints and request/response contracts, handles duplicate requests / failed disbursements / partial failures, and lets the caller track status after submission.

WHAT I CHOSE TO DEMONSTRATE

The **Pull-and-Pay funding model with a compliance gate** — the most complex realistic flow. It forces the full design: async compliance, explicit fund pull state, ambiguous-failure handling, and a refund path. More representative of a real platform than the simpler prefunded model.

4 Assumptions

Each resolves a scenario the task leaves open.

#	Assumption	Key trade-off
1	Submission sync, settlement async	202 Accepted + async tracking vs. coupling API uptime to rail latency
2	Single payout = primary unit; batch is additive	Get idempotency right on one payout first, then compose
3	Platform picks the rail, caller doesn't	Hides internal topology; Boku's value prop is the 65-country coverage
4	Pull-and-Pay as primary showcase	Forces full design: compliance gate, fund pull state, refund flow

Assumption 1 — Async by Design

Scenario: `POST /payouts` hits the API. The actual funds movement can take seconds to multiple business days depending on corridor and rail.

Decision: `202 Accepted` immediately. Acceptance path = thin DB write + fire-and-forget Kafka publish, **< 50ms P95** — comfortably inside any API Gateway timeout.

Alternative rejected: Fully synchronous confirm-then-respond — couples API availability to partner rail latency. One slow rail makes the entire API appear down.

```
POST /payouts      → 202 Accepted { payout_id, status: PENDING }
                    ↓
                    async pipeline
                    ↓
GET /payouts/{id}   → { status: SETTLED, history: [...] }
```

Assumption 3 — Rail Selection Is the Platform's Job

Scenario: Boku is global (65 countries) — Singapore, Kenya, and the US ride completely different rails.

Decision: Caller specifies **what** and **to whom**. Platform decides **how**.

Caller-selects-rail	Platform-selects-rail (chosen)
Explicit "rail": "LOCAL_BANK_SG" in request	amount , currency , beneficiary only
Leaks internal partner topology into public contract	Rail selection + failover internal
Couples API to infra decisions	Caller API stays simple

Open question Q1: Is Boku a managed service (platform picks rail) or a connectivity layer (caller picks rail)? If the latter, the design simplifies — rail selection layer disappears entirely.

Assumption 4 — Pull-and-Pay as the Showcase

Model	How it works	Design impact
Prefunded wallet	Merchant tops up Boku balance in advance	Synchronous <code>INSUFFICIENT_FUNDS</code> . Simpler — no refund flow.
Pull-and-Pay (primary)	Platform debits caller via direct debit mandate as a dedicated async step	Two failure surfaces. Full design: compliance gate, <code>FUND_PULLING</code> state, refund flow.

Why Pull-and-Pay forces the right design:

- Compliance gate must run **before** any funds move
- `FUND_PULLING` as a separate state means reaching `SUBMITTED` **guarantees** funds are with Boku — no `funds_pulled` flag needed anywhere
- Every failure from `SUBMITTED` triggers a refund unconditionally

Open question Q3: Which model does Boku actually use? If prefunded, the design simplifies: remove `FUND_PULLING` , `FUND_PULL_FAILED` , and the refund flow.

API Surface

Endpoint	Purpose
POST /payouts	Submit a single payout — Idempotency-Key required
GET /payouts/{id}	Current status + full lifecycle history — always the source of truth
GET /payouts?external_ref={ref}	Recovery path when caller crashed before persisting payout_id
POST /payouts/batch	Fan-out to N independent payout records
GET /payouts/batch/{batchId}	Rollup: succeeded/failed/pending counts
POST /payouts/{id}/cancel	State-machine guard — only PENDING is cancellable
POST /webhooks/subscriptions	Register callback URL — push complement to polling
POST /quotes	Lock FX rate before payout — optional, short TTL
GET /quotes/{quoteId}	Check quote validity

DELETE /payouts/{id} is absent — funds-movement records are never deleted, only transitioned to terminal states.

Payout State Machine

PENDING → PENDING_COMPLIANCE → FUND_PULLING → SUBMITTED → SETTLED	← happy path
PENDING → FUND_PULLING → SUBMITTED → SETTLED	← no compliance gate
PENDING → SUBMITTED → SETTLED	← prefunded wallet
PENDING_COMPLIANCE → REJECTED_COMPLIANCE	← hard AML hit (terminal, no funds moved)
PENDING_COMPLIANCE → PENDING_MANUAL_REVIEW	← soft hit (human review)
PENDING_MANUAL_REVIEW → FUND_PULLING REJECTED_COMPLIANCE	← reviewer decides
FUND_PULLING → SUBMITTED	← pull succeeds, funds with Boku
FUND_PULLING → FUND_PULL_FAILED	← terminal, no refund needed
SUBMITTED → SETTLED FAILED RETURNED SUBMITTED_UNCONFIRMED	
SUBMITTED_UNCONFIRMED → SETTLED FAILED	← partner-ref lookup only
FAILED → REFUND_PENDING → REFUNDED REFUND_FAILED	
RETURNED → REFUND_PENDING → REFUNDED REFUND_FAILED	
PENDING → CANCELLED	

Two-layer write guard — not just one:

`UPDATE ... WHERE status = 'X'` handles duplicate Kafka events on the same worker. But it doesn't cover **two independent writer types** reading the same row simultaneously before either commits — e.g. a partner callback + midnight batch reconciliation both seeing `SUBMITTED` and both writing `SETTLED`.

Optimistic locking with a `version` column closes this gap:

Compliance Gate — Before Any Funds Move

```
payout.created (Kafka)
```

```
↓
```

```
Fraud screening      ← velocity, amount anomaly, beneficiary pattern
```

```
PASS → AML scan      ← OFAC, UN, EU, local sanctions + PEP lists
```

```
PASS → payout.compliance_passed → FUND_PULLING
```

```
HARD HIT → REJECTED_COMPLIANCE (terminal, auto, may trigger regulatory reporting)
```

```
SOFT HIT → PENDING_MANUAL_REVIEW (human review, SLA-bound)
```

```
FAIL → REJECTED_COMPLIANCE (terminal)
```

Why `REJECTED_COMPLIANCE` ≠ `FAILED`:

- `FAILED` = rail issue, potentially retryable with a different account
- `REJECTED_COMPLIANCE` = legal/policy block — no retry changes the outcome; may trigger reporting obligation

Why compliance runs before fund pull:

No money moves until the payout is clean. `REJECTED_COMPLIANCE` requires no refund — nothing was debited.

The SUBMITTED_UNCONFIRMED Problem

Scenario: Disbursement worker sends to Partner Rail. Connection drops. Did the partner process it?

```
DW → Rail: submit disbursement (client_reference = Idempotency-Key)
Rail: [timeout / network error]
DW → DB: UPDATE status = SUBMITTED_UNCONFIRMED
```

The wrong responses:

- ✗ Trigger refund immediately → refund money that was already disbursed
- ✗ Failover to Partner B → **double payment** if Partner A processed it

The right response:

```
DW → Rail: GET /status?ref=<Idempotency-Key>
    SETTLED → update SETTLED, notify merchant
    NOT FOUND / FAILED → update FAILED → trigger refund
```

`SUBMITTED_UNCONFIRMED` is deliberately non-terminal. Failover is blocked until the partner-reference lookup resolves the outcome.

Refund Trigger — Simplified by FUND_PULLING Separation

State	Refund needed?	Why
REJECTED_COMPLIANCE	No	Compliance runs before any funds move
FUND_PULL_FAILED	No	Caller's account was never debited
FAILED from SUBMITTED	Yes — always	SUBMITTED = funds are with Boku
RETURNED from SUBMITTED	Yes — always	Rail reversed after Boku held funds
SUBMITTED_UNCONFIRMED → FAILED	Yes — after lookup	Do not refund while outcome unknown

The FUND_PULLING insight:


Before: need a `funds_pulled` boolean, checked in the refund worker.

After: reaching `SUBMITTED` is the guarantee. Refund worker does exactly one thing.

REFUND_FAILED — money is with Boku, held by neither party. Zero-tolerance alert: page oncall immediately, no acknowledgement window.

Refund Design — REFUND_REQUIRED + Batch Job

Updated state machine for refunds:



The trade-off — User Satisfaction vs System Performance:

Dimension	Immediate (event-driven)	Batch Job via REFUND_REQUIRED (chosen)
User satisfaction	Faster refund — merchant gets money back sooner	Delay of up to one batch interval (e.g. 5 min)
System risk	Burst of failures = burst of refund rail calls	Controlled rate — batch job throttles calls
Observability	State scattered across Kafka consumer metrics	SELECT WHERE status='REFUND_REQUIRED' — everything visible
Failure recovery	Kafka DLQ — harder to inspect and replay	DB rows are durable, queryable, trivially reprocessable
Concurrency	Consumer race on same payout	SELECT FOR UPDATE SKIP LOCKED — safe, no double-execution

Safe batch pickup:

```
SELECT payout_id, version FROM payouts
WHERE status = 'REFUND_REQUIRED' ORDER BY updated_at ASC LIMIT 100
FOR UPDATE SKIP LOCKED;
-- advance to REFUND_PENDING before the rail call
```

Cancellation — Boundary at the First External Call

Rule: `POST /payouts/{id}/cancel` is only permitted **before any external API call has been made**.

Status	Cancellable?	Why
PENDING	✔ Yes	DB insert + Kafka only — no external call
PENDING_COMPLIANCE	✔ Yes	Compliance is Boku-internal — no partner contacted
PENDING_MANUAL_REVIEW	✔ Yes	Human queue — safe to withdraw
FUND_PULLING	✘ No	Fund pull API already called — unknown mid-flight state
SUBMITTED and beyond	✘ No	Disbursement rail called — use refund flow instead
Any terminal state	✘ No	Already resolved — return 422

Response on non-cancellable state — 422 , not 409 :

```
{ "error_code": "CANCELLATION_NOT_ALLOWED",  
  "message": "Payout is in FUND_PULLING. Cancellation not permitted after external API call.  
             Funds will be refunded automatically if disbursement fails." }
```

409 implies a retry might work. 422 is unambiguous: the window is closed.

Race condition — cancel vs. compliance-pass arriving simultaneously:

```
UPDATE payouts SET status='CANCELLED', version=version+1
```

Idempotency — Three Layers, One Principle

Idempotency-Key vs. Correlation-ID — don't collapse them:

- `Idempotency-Key` — caller-supplied, stable across all retries. *"Is this a duplicate?"*
- `Correlation-ID` — generated per attempt, threaded through logs/events/spans. *"How do I trace this execution?"*

THREE ENFORCEMENT POINTS

Layer	Mechanism
Ingress (caller → API)	<code>Idempotency-Key</code> header + Postgres unique constraint in same ACID transaction as INSERT
Internal (Kafka consumers)	<code>UPDATE ... WHERE status = 'X'</code> — duplicate event → no-op, not double transition
Egress (service → partner rail)	<code>client_reference = Idempotency-Key</code> on every outbound call — enables <code>SUBMITTED_UNCONFIRMED</code> recovery

Also covers the 30s API Gateway timeout: caller retries same key → dedup returns original response. No new state or logic needed.

Observability — Audit Table as Foundation

```
CREATE TABLE payout_audit (  
  id          BIGSERIAL PRIMARY KEY,  
  payout_id   VARCHAR(40) NOT NULL,  
  correlation_id VARCHAR(40) NOT NULL,  
  from_status VARCHAR(40),           -- NULL for first PENDING row  
  to_status   VARCHAR(40) NOT NULL,  
  triggered_by VARCHAR(60) NOT NULL, -- 'api', 'compliance-service', 'fund-pull-worker'...  
  duration_ms INTEGER,              -- time spent IN this step  
  failure_reason VARCHAR(80),  
  partner_ref   VARCHAR(120),  
  created_at    TIMESTAMPTZ NOT NULL DEFAULT NOW()  
);
```

Append-only. One row per state transition. Never updated. The mutable `status` column is a projection — the audit table is the immutable ledger.

What it unlocks: per-step SLA monitoring, P95 latency dashboards, proactive alerting on step duration, full trace of every payout's lifecycle.

Top 5 Alerts (from 18 defined)

Priority	Alert	Condition	Why selected
☑	POST_PAYOUTS_CRITICAL	P95 response > 1,000ms for 1min	Only step where every caller is affected simultaneously
☑	FAIL_RATE_CRITICAL	> 15% submissions failed in 5min	Systemic rail problem already hitting real volume — every minute compounds
☑	REFUND_FAILED_ANY	Any REFUND_FAILED in 5min	Zero-tolerance — money with Boku, no ack window, page immediately
☑	SUBMITTED_UNCONFIRMED_AGE	Any payout in SUBMITTED_UNCONFIRMED > 15min	Silent double-payment risk if failover triggers without resolution
☑	REFUND_PENDING_AGE	Any payout in REFUND_PENDING > 30min	Caller's money held with no progress — regulatory issue in some jurisdictions

15min and 30min thresholds are **policy decisions**, not technical ones. Good to surface live — what is Boku's tolerance?

AWS Architecture

Component	Service	Why
Public API entry	API Gateway	Auth, rate limiting, request validation
Payout services	EKS	6 Spring Boot microservices: Payout Service, Compliance Service, Fund Pull Worker, Disbursement Worker, Refund Worker, Webhook Dispatcher
Event backbone	MSK (Kafka)	6 topics — durable, ordered, replayable
System of record	RDS PostgreSQL	ACID for payout ledger + <code>payout_audit</code> + idempotency-key table
Retry / DLQ	SQS + DLQ	Per-rail exponential backoff; exhausted messages for follow-up
Idempotency fast-path	ElastiCache (Redis)	Sub-ms duplicate-key check before hitting Postgres
Scheduled reconciliation	Lambda + EventBridge	Matches partner reports against <code>payout_audit</code>
Observability	CloudWatch + X-Ray	Per-step duration metrics, distributed traces keyed on Correlation ID

Language/framework: Java 21 + Spring Boot (matches Boku stack). WebFlux for Disbursement Worker and Fund Pull Worker — high fan-out to slow rails is the non-blocking use case it's built for.

Stakeholder Coverage — Honest Assessment

Role	Covered	Gap
Customer	Idempotency, dual pull+push tracking, structured error codes, full history	No stated SLA on <code>PENDING_COMPLIANCE</code> wait time
Developer	Full state machine, Correlation-ID vs Idempotency-Key split, consistent error shape	No test strategy, no API versioning scheme
DevOps	Full AWS mapping, 6 EKS services, MSK/RDS/SQS/Secrets Manager	No canary/blue-green strategy, no autoscaling policy
Operation Team	Audit table, 5 priority alerts with thresholds, DLQ for retries	No admin API; runbooks are stubs — data supports them, procedures not written
Financial Team	Append-only audit trail, reconciliation design, REFUNDED trail	No GL integration point, no retention policy
Stakeholder	Greenfield rationale, AWS choices with precedent	No TCO estimate, no phased timeline

Strongest on: Customer / Developer / DevOps — direct extensions of DASH-proven patterns.

Thinnest on: Financial Team / Stakeholder — name these gaps first, don't wait to be caught.

Open Questions for the Walkthrough

Q1 — Managed service or connectivity layer?

If caller specifies the rail → rail selection layer gone, failover-blocking in `SUBMITTED_UNCONFIRMED` simplified. Changes the entire §3.9 story.

Q2 — If platform picks rail: priority-ordered fallback or multi-rail scoring?

- Priority-ordered fallback → circuit breaker model, current design applies
- Multi-rail scoring → need a selection step not currently designed

Q3 — Prefunded wallet or Pull-and-Pay?

If prefunded → remove `FUND_PULLING` / `FUND_PULL_FAILED` / refund flow; add wallet debit in acceptance transaction; synchronous `INSUFFICIENT_FUNDS`. Show you know both.

Demo

Live API Playground

[https://boku-payouts-design.vercel.app/api-docs/
?server=https://mywiremockserver-production.up.railway.app](https://boku-payouts-design.vercel.app/api-docs/?server=https://mywiremockserver-production.up.railway.app)

DEMO IDS TO TRY

- `pyo_SETTLED_001` — happy path with full history + timing
- `pyo_FAILED_001` — `BENEFICIARY_ACCOUNT_INVALID`
- `pyo_UNCONF_001` — `SUBMITTED_UNCONFIRMED`
- `btc_DEMO_001` — batch: 8 settled / 2 failed
- `qte_EXPIRED_001` — expired FX quote